

Building a content-based recommender system for scientific papers

Capstone Project – Springboard Data Science Career Track

Hehuan Liao

Oct. 8, 2020

Problem Identification

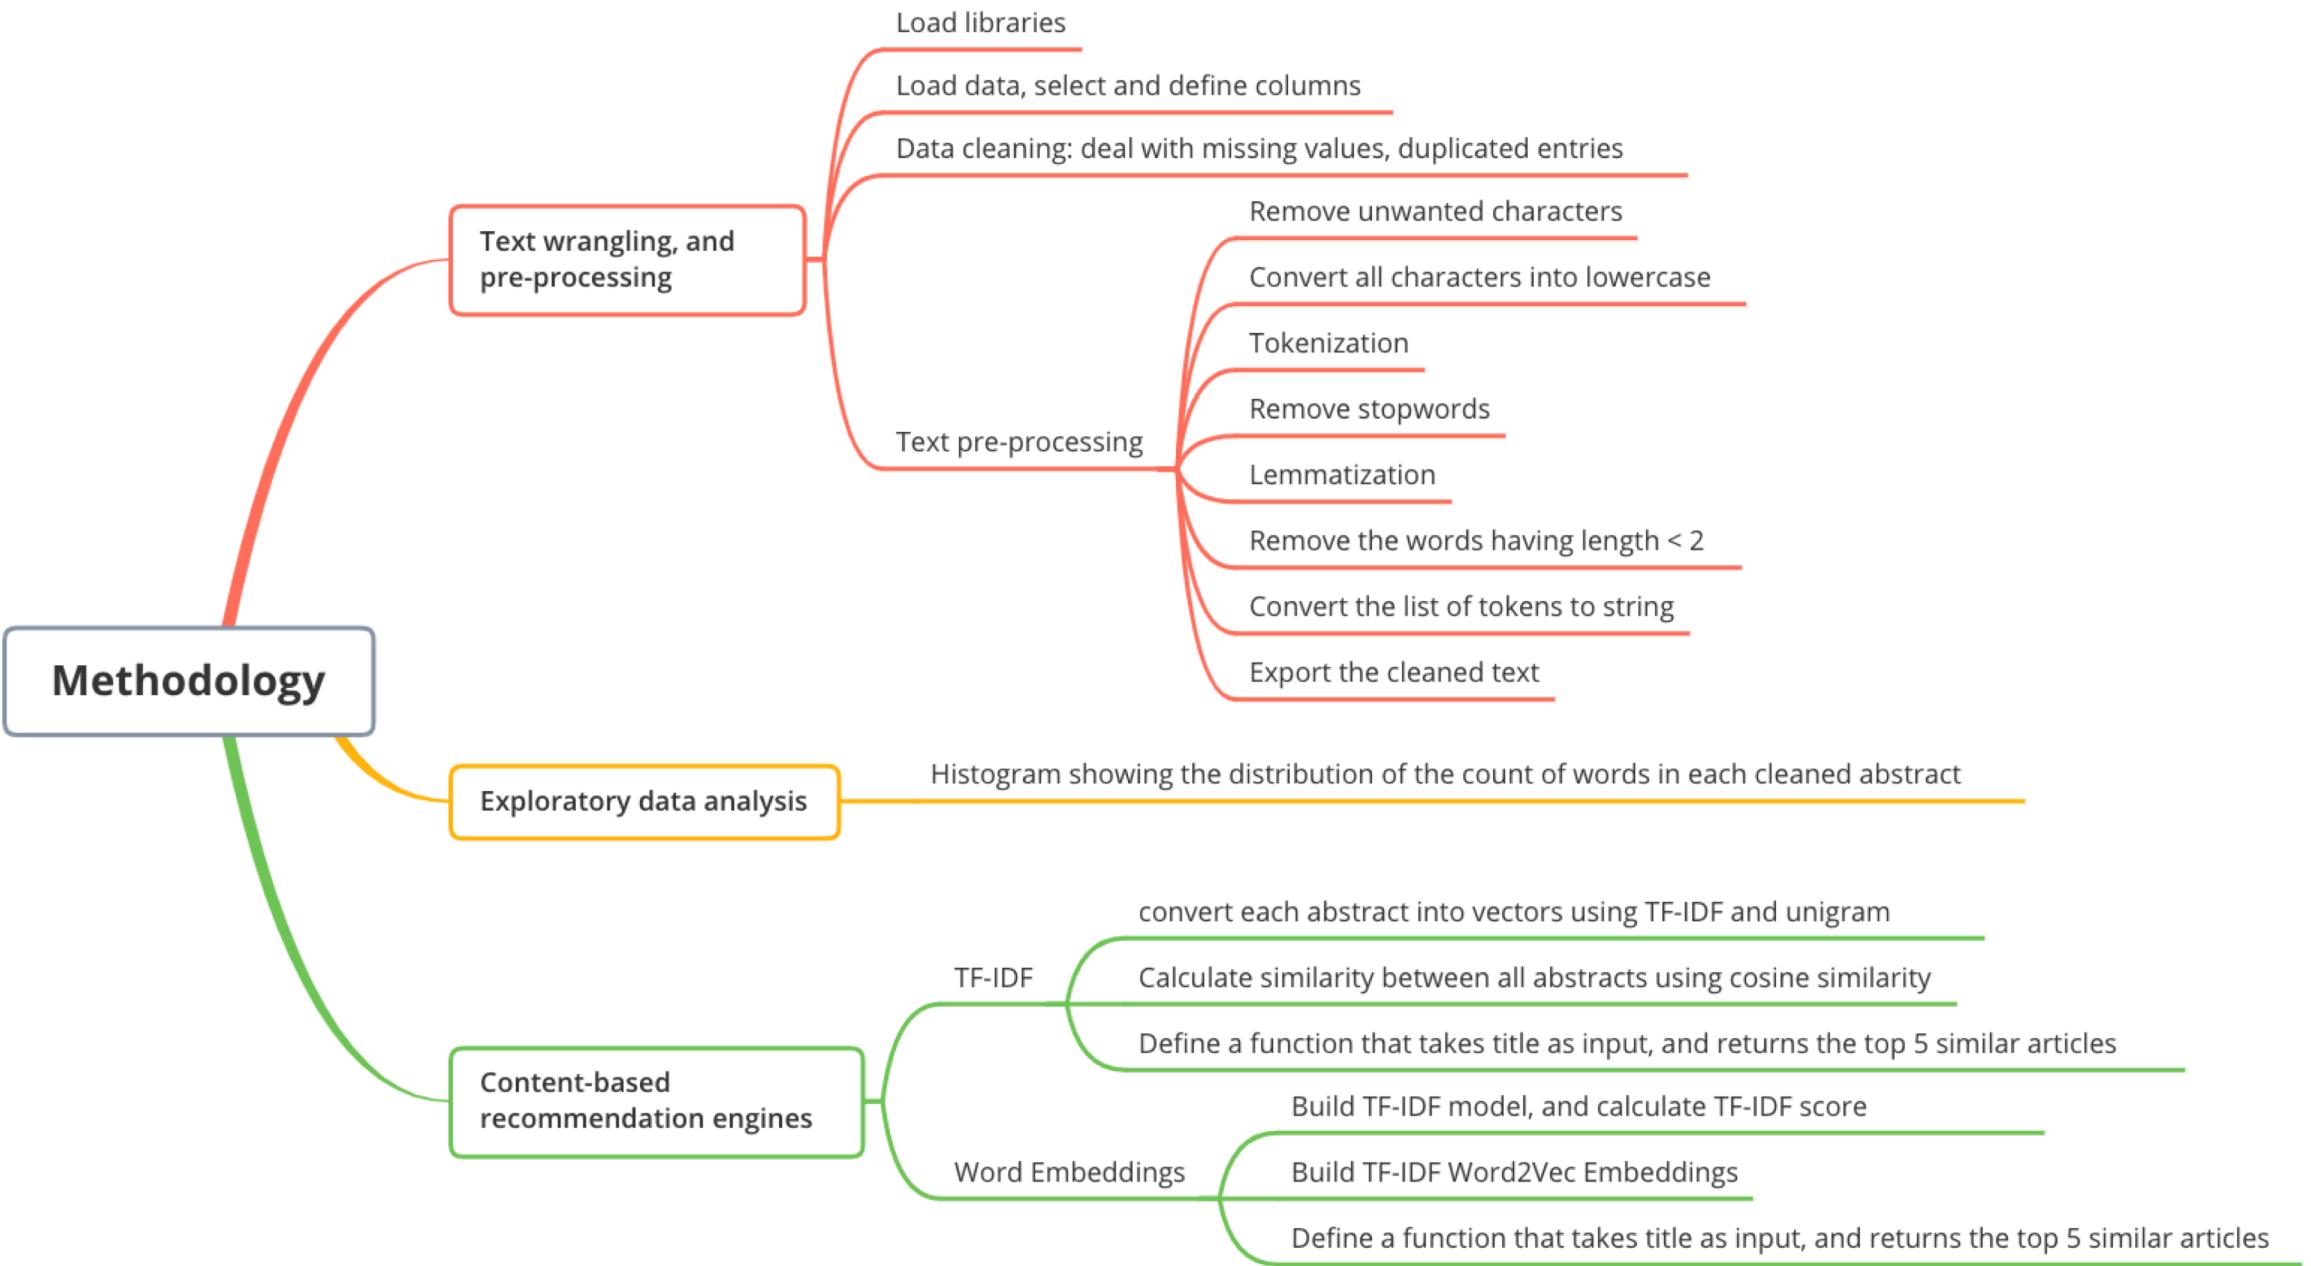
- ArXiv is an open-access repository of e-prints of scientific papers covering various fields
- As of today, arXiv has hosted over 1.7 million scholarly articles, which is increasing dramatically
- Extracting relevant information efficiently becomes a challenge
- A recommender system would make the exploration of scholarly articles faster and more accurate

👉 Overall goal: develop a recommender system that can present the users with related articles based on the article they are reading.

Dataset description

- Data source: arXiv dataset from Kaggle
- Over 1.7 million observations of 10 variables

No.	Variable name	Description
1	id	ArXiv ID
2	submitter	Who submitted the paper
3	authors	Authors of the paper
4	title	Title of the paper
5	comments	Additional info, such as number of pages and figures
6	Journal-ref	Information about the journal the paper was published in
7	doi	Digital object identifier
8	abstract	The abstract of the paper
9	categories	Categories / tags in the ArXiv system
10	versions	A version history



Text wrangling and pre-processing

```
# abstract before pre-processing: abstract  
dat['abstract'][0]
```

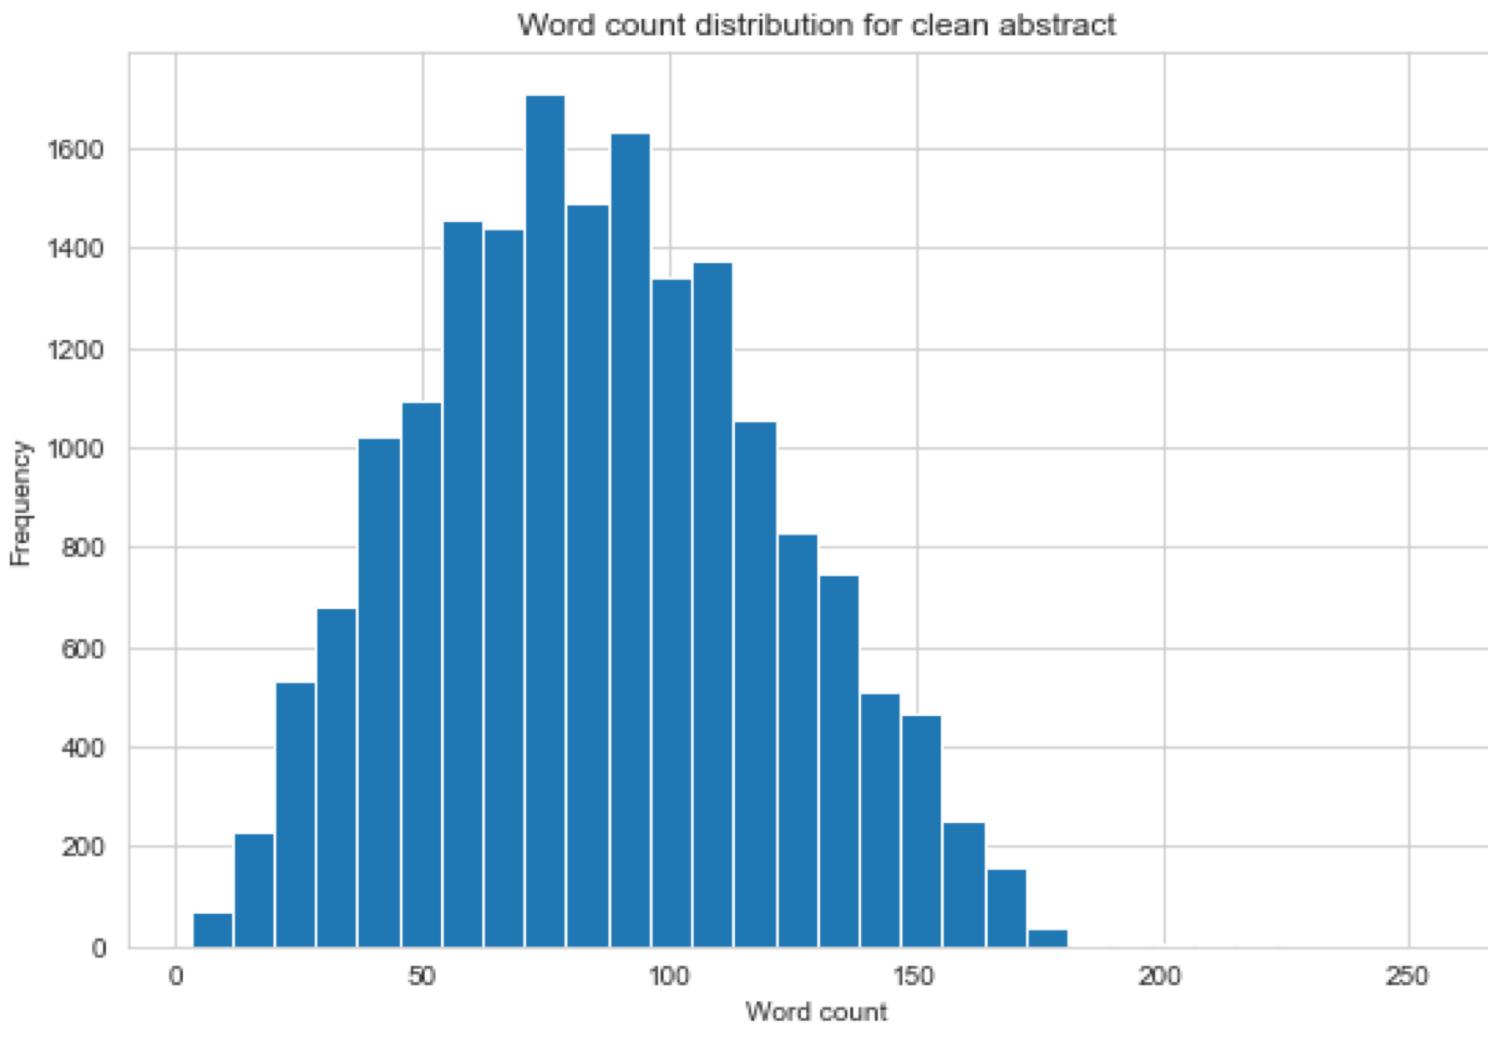
" We present and study a family of metrics on the space of compact subsets of \mathbb{R}^N (that we call ``shapes''). These metrics are ``geometric'', that is, they are independent of rotation and translation; and these metrics enjoy many interesting properties, as, for example, the existence of minimal geodesics. We view our space of shapes as a subset of Banach (or Hilbert) manifolds: so we can define a ``tangent manifold'' to shapes, and (in a very weak form) talk of a ``Riemannian Geometry'' of shapes. Some of the metrics that we propose are topologically equivalent to the Hausdorff metric; but at the same time, they are more ``regular'', since we can hope for a local uniqueness of minimal geodesics.\n We also study properties of the metrics obtained by isometrically identifying a generic metric space with a subset of a Banach space to obtain a rigidity result.\n"

```
# abstract after pre-processing: clean_abstract  
dat['clean_abstract'][0]
```

'present study family metrics space compact subsets call shape metrics geometric independent rotation translation metrics enjoy many interest properties example existence minimal geodesics view space shape subset banach hilbert manifold define tangent manifold shape weak form talk riemannian geometry shape metrics propose topologically equivalent hausdorff metric time regular since hope local uniqueness minimal geodesics also study properties metrics obtain isometrically identify generic metric space subset banach space obtain rigidity result'

A comparison of an example abstract before and after preprocessing

Exploratory data analysis (EDA)



count	18118.000000
mean	85.373938
std	35.237470
min	3.000000
25%	59.000000
50%	84.000000
75%	110.000000
max	257.000000

Content-based recommendation engines

- **TF-IDF (term frequency – inverse document frequency)**
 - quantify a word in documents that is weighted to signify its importance in the document and corpus
 - $\text{TF-IDF} = \text{TF}(w) * \text{IDF}(w)$
 $\text{TF}(w) = (\# \text{ of times term } w \text{ appears in a document}) / (\text{total } \# \text{ of terms in the document})$
 $\text{IDF}(w) = \log_e(\text{Total } \# \text{ of documents} / \# \text{ of documents with term } w \text{ in it})$
- **Word embeddings produced via Word2Vec algorithm:**
 - Word embedding is dense representation of words in the form of numeric vectors
 - Word2vec algorithm uses a neural network model to learn word associations from a large corpus of text

Recommendation engine: TF-IDF

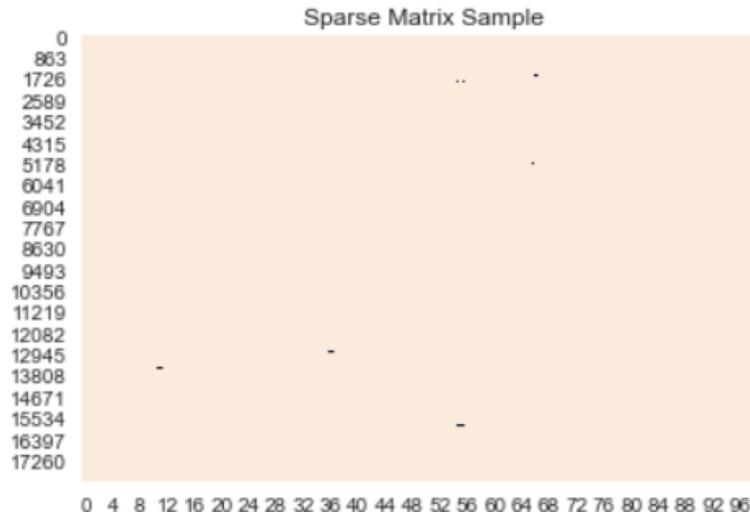
```
: from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity

#convert the abstract into vectors and use unigram
tfidf = TfidfVectorizer()
tfidf_matrix = tfidf.fit_transform(dat['clean_abstract'])
df1 = pd.DataFrame(tfidf_matrix.todense(),columns=tfidf.get_feature_names())
df1.shape

: (18118, 37948)

: #plot examples of tfidf_matrix
sns.heatmap(tfidf_matrix.todense()
             [:,np.random.randint(0,10000,100)]==0, vmin=0, vmax=1,
             cbar=False).set_title('Sparse Matrix Sample')

: Text(0.5, 1.0, 'Sparse Matrix Sample')
```



TF-IDF creates a sparse, high-dimensional feature, and does not capture the semantic meaning

Recommendation engine: TF-IDF

```
paper_recommend1('Banach-like metrics and metrics of compact sets')
```

Here are the top 5 similar articles that may interest you:

#1. Title: "Efficient PDE constrained shape optimization based on Steklov-Poincar\'e type metrics". (DOI:10.1137/15M1029369)

#2. Title: "A Locally Weighted Fixation Density-Based Metric for Assessing the Quality of Visual Saliency Predictions". (DOI:10.1109/TIP.2016.2577498)

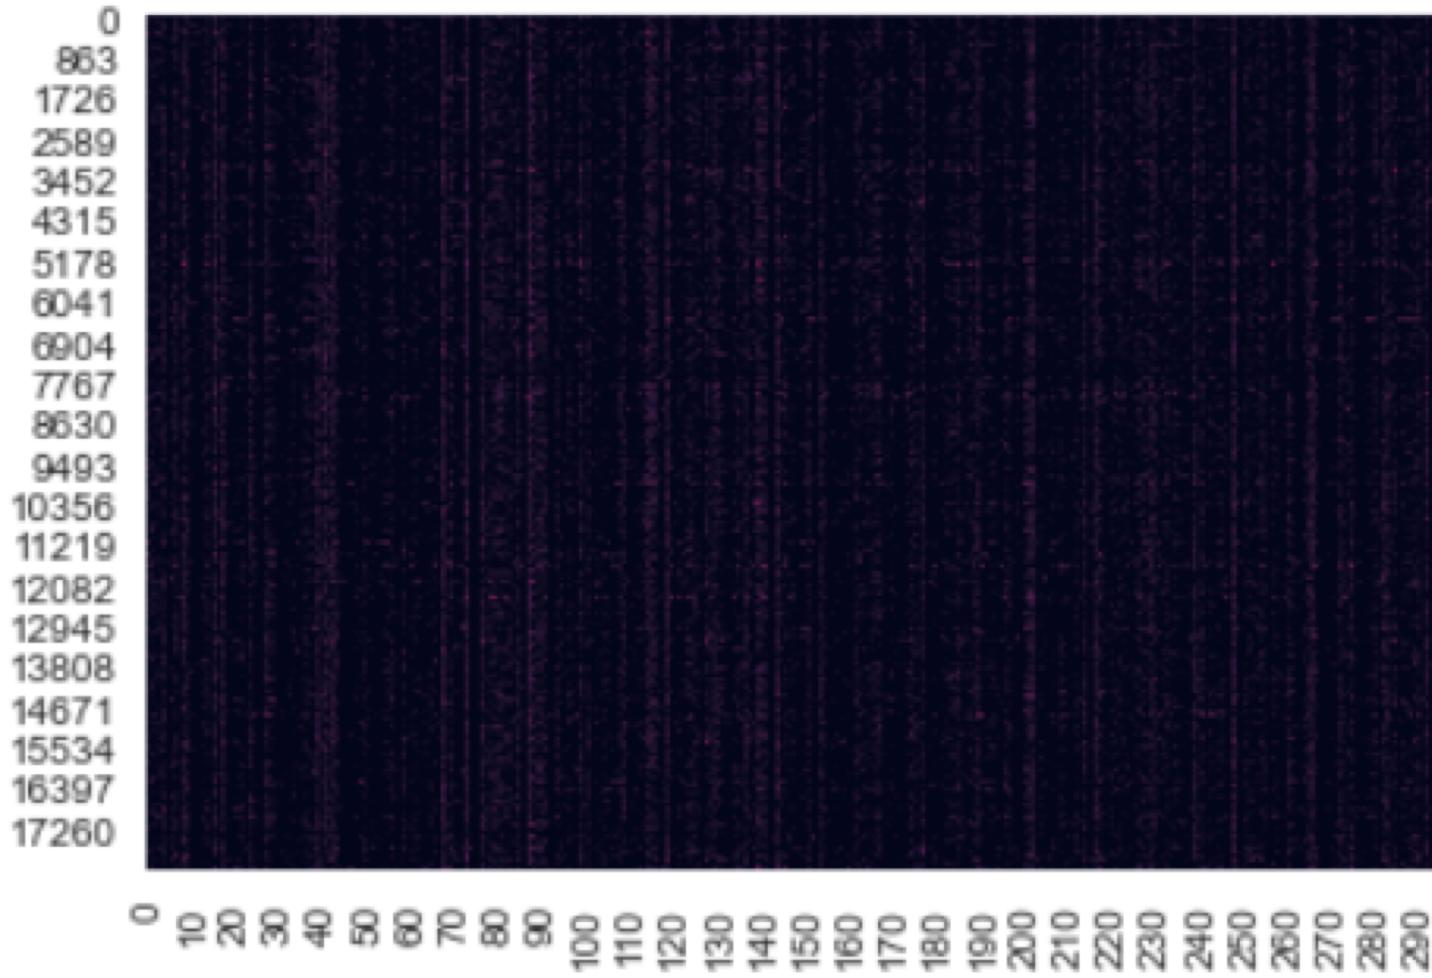
#3. Title: "Sparse Geometric Representation Through Local Shape Probing". (DOI:10.1109/TVCG.2017.2719024)

#4. Title: "Sphere equivalence, Banach expanders, and extrapolation". (DOI:10.1093/imrn/rnu075)

#5. Title: "Hexagonal geodesic 3-webs". (DOI:None)

Recommendation engine: word embeddings

Dense Matrix



Word2Vec creates a dense, low dimensional feature, and captures the semantic meaning

Recommendation engine: word embeddings

```
paper_recommend2('Budget Constraints in Prediction Markets')
```

Here are the top 5 similar articles that may interest you:

- #1. Title: "Modeling the residential electricity consumption within a restructured power market". (DOI:None)
- #2. Title: "Phase Transition in the S&P Stock Market". (DOI:10.1007/s11403-015-0160-x)
- #3. Title: "Investigating the effect of competitiveness power in estimating the average weighted price in electricity market". (DOI:10.1016/j.tej.2019.106628)
- #4. Title: "Hierarchical structure of stock price fluctuations in financial markets". (DOI:10.1088/1742-5468/2012/12/P12016)
- #5. Title: "Satiation in Fisher Markets and Approximation of Nash Social Welfare". (DOI:None)

Conclusions

With the exponentially increasing number of scholarly articles, extracting relevant information efficiently becomes a challenge. Recommendation systems provide a means to make the exploration of scholarly articles faster and more accurately, which would of great importance.

This project explored the use of basic TF-IDF to build a recommendation engine that takes the title of an article that a user is reading, and returns top 5 related articles that may also interest the user. However, TF-IDF does not capture the semantic meaning, and creates a sparse, high-dimensional feature. Word embeddings create a dense, low-dimensional feature, and captures the semantic meaning quite well. Therefore, we explored the use of TF-IDF Word2Vec to build a similar recommendation engine.

Both of the two content-based recommendation engines were able to recommend the top 5 related articles to the article that a user was reading. However, as no user-based data was currently available, we cannot evaluate the model performance quantitatively. Future improvements would include additional features such as authors, categories of articles, and user-based data.

Acknowledgements

- Data source: arXiv database
- Springboard mentor Dipanjan Sarkar, and the Springboard team

The entire project is available:

https://nbviewer.jupyter.org/github/hehuanliao/Springboard/blob/master/Capstone3/article_recommend.ipynb